Productionizing

AI + LLM Apps with Ray Serve

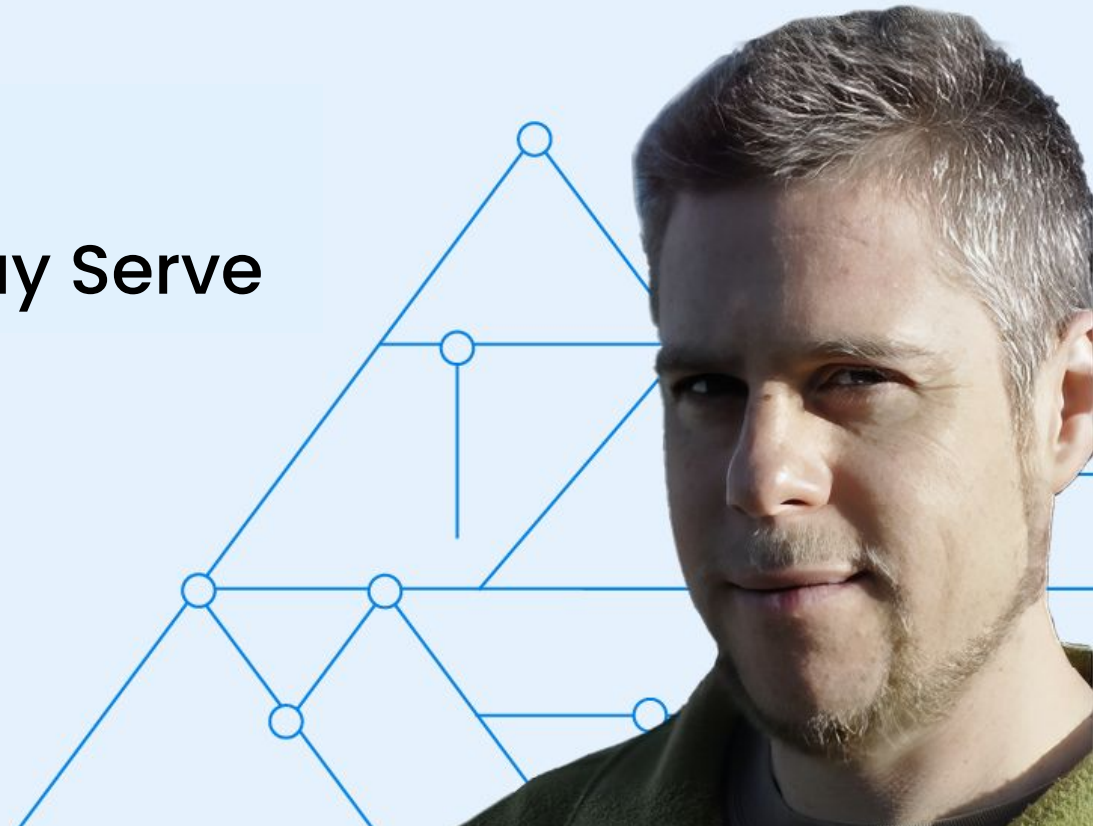

Adam Breindel
Anyscale

adamb@anyscale.com
@adbreind

🧑‍💻 Meet the tutorial team!

Marwan
marwan@anyscale.com

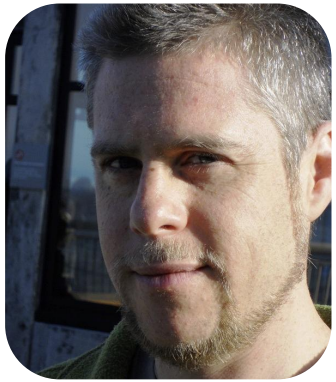

Adam
adamb@anyscale.com

Kamil
kamil@anyscale.com

The Plan

Here’s what to expect today.

Today’s agenda.

- What is Ray Serve?
- Why use Ray and Ray Serve for scalable AI?
- Build complex ML applications with Ray + Serve
- Under the hood: features for powering production apps
- Architecture options, hands-on labs, and Q&A

✅ Tech check.

 Accessing Anyscale clusters.

- All work will be in Anyscale provisioned clusters.
- Our GitHub repo will be mounted automatically.
- Access begins now.
  - Check your email for login information.
  - Step-by-step instructions to follow.

Anyscale login

Link to Anyscale cluster: console.anyscale.com

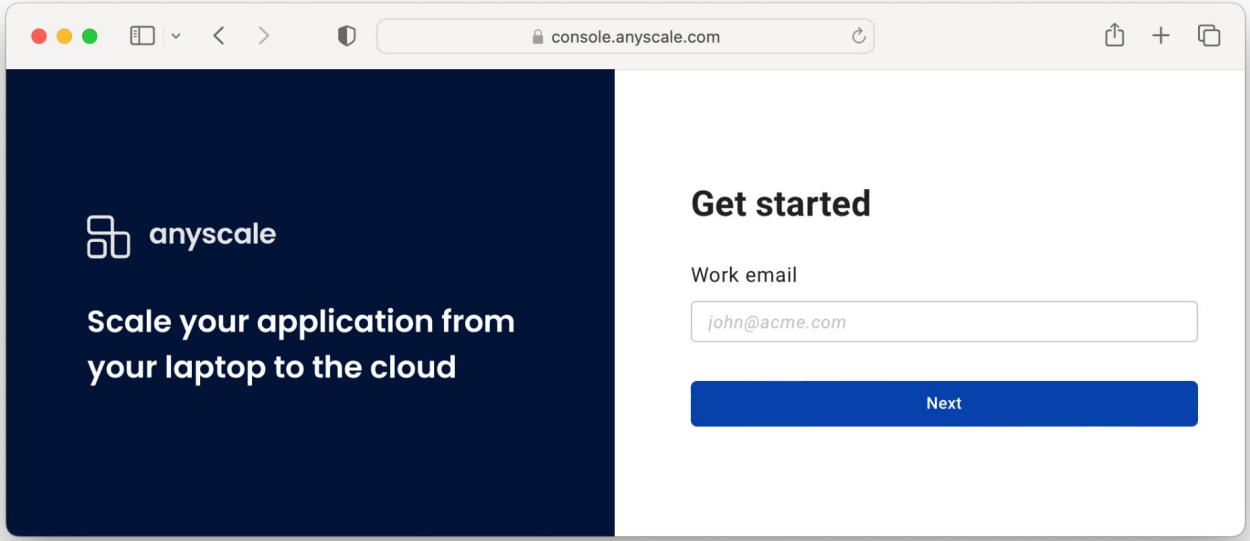

Enter the **unique credentials** sent to your email!

1. Select Workspaces

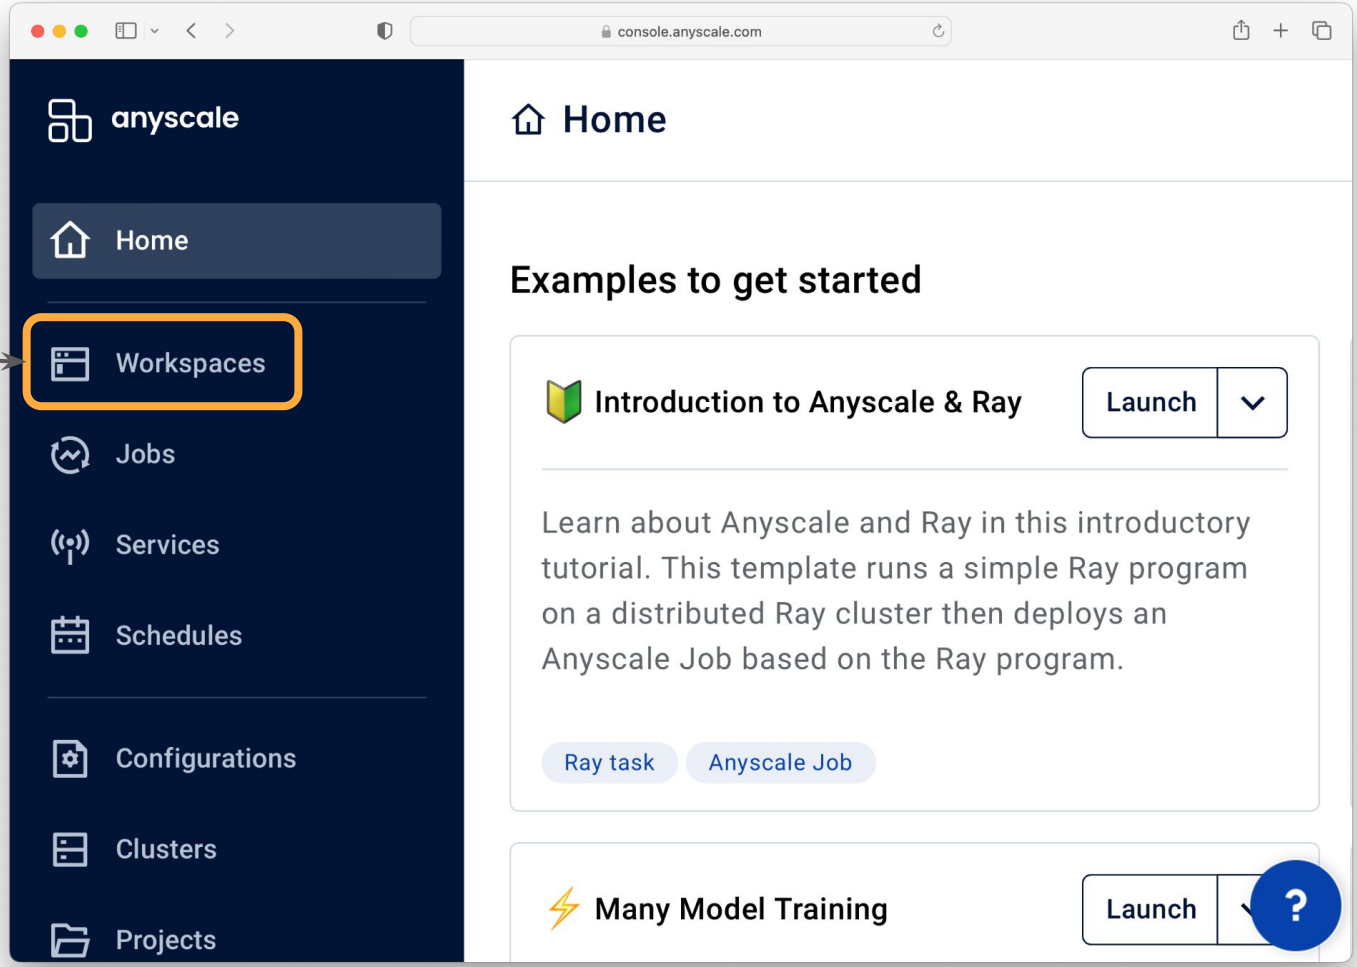

2. Select Your Workspace

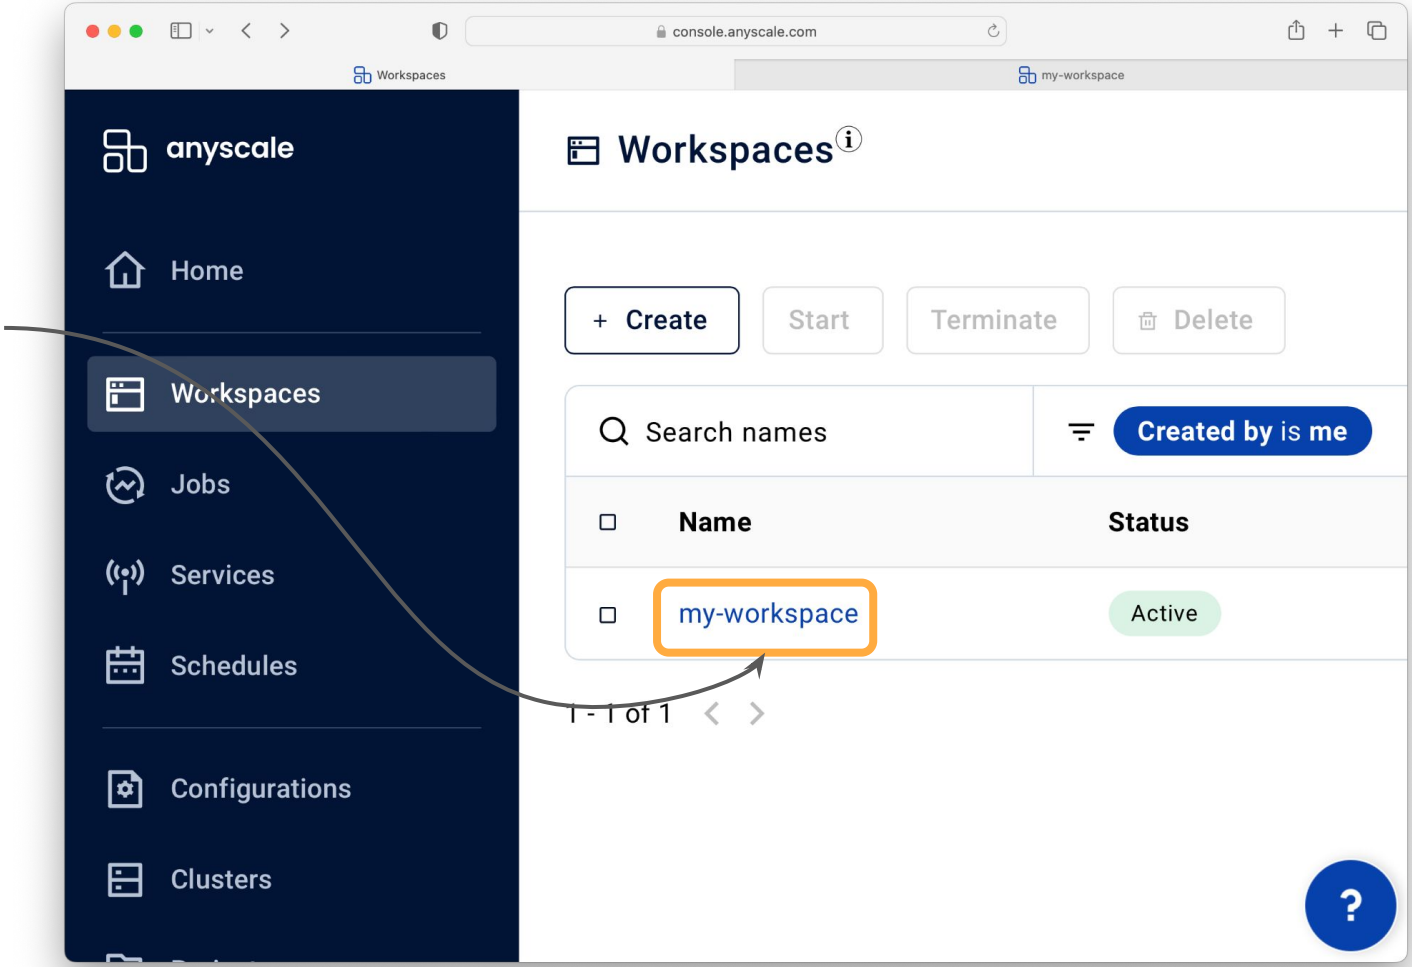

3. Click on Jupyter icon

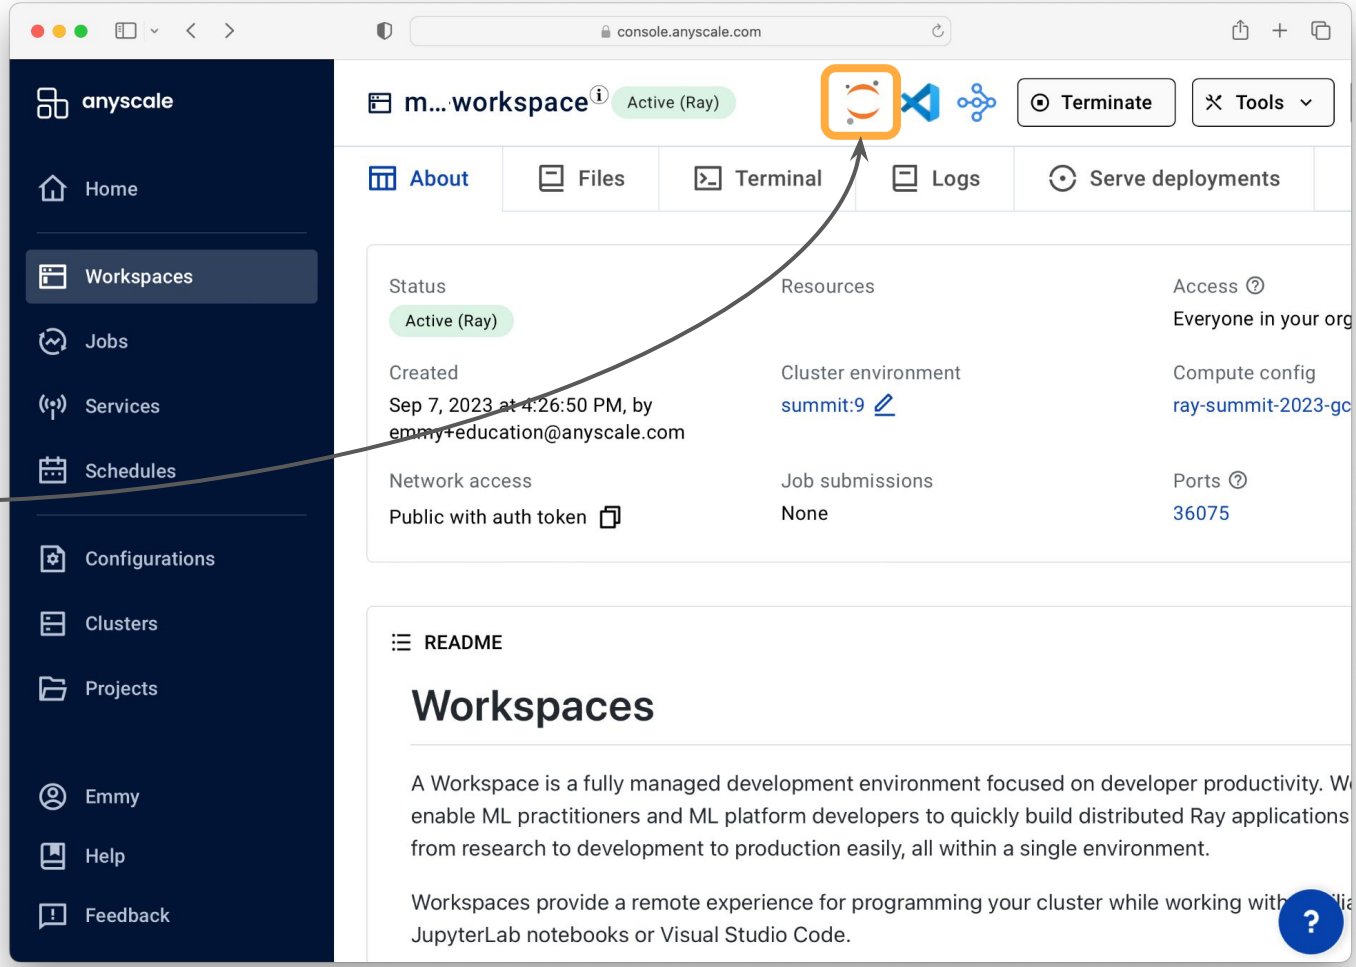

4. Find the content for your class here.

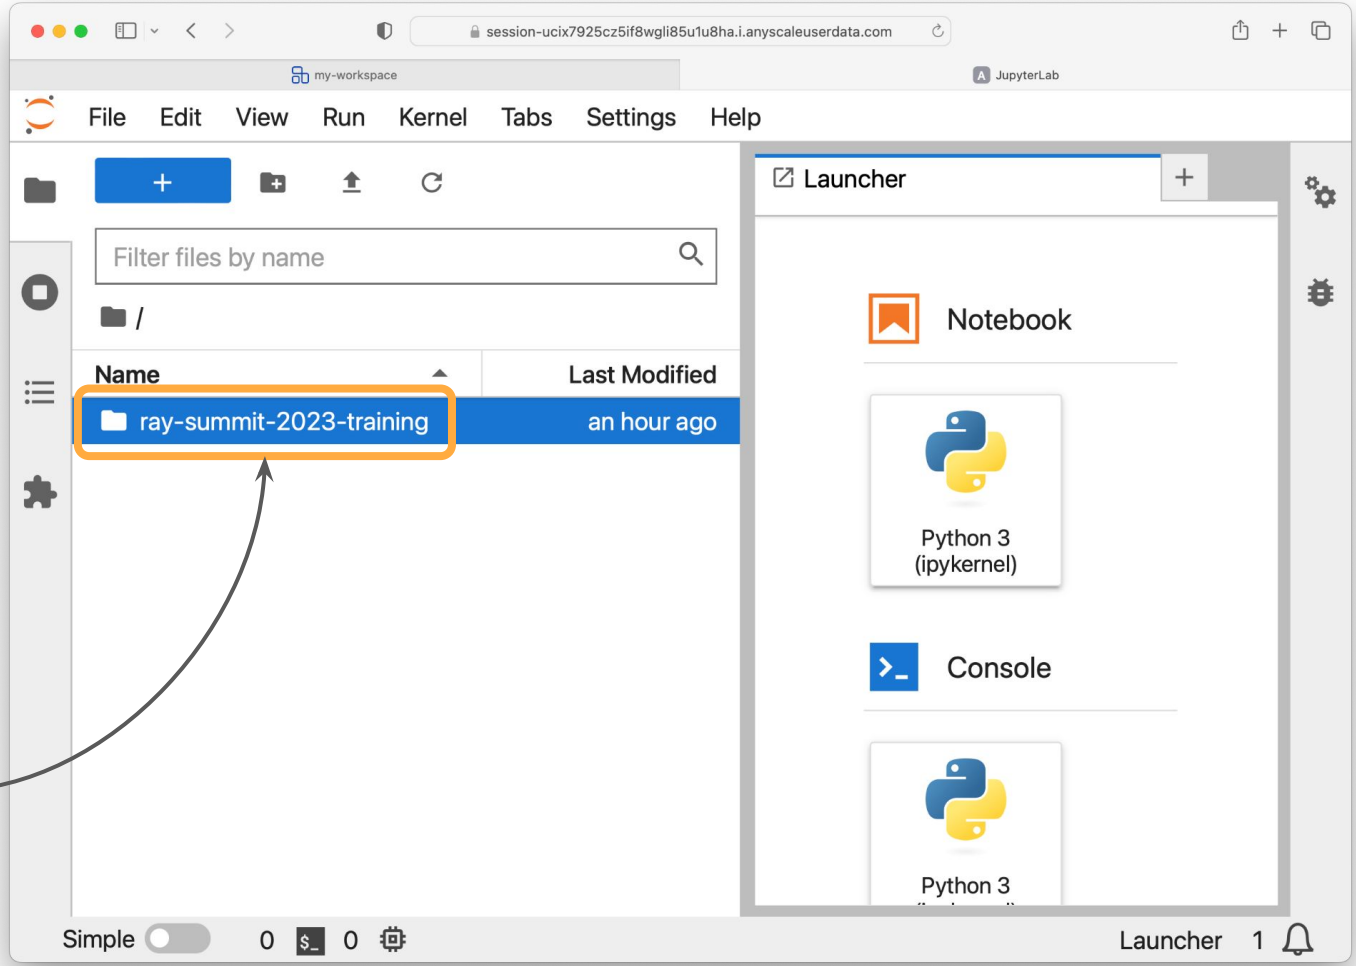

Time for a Break!

15 minutes.

🍎 Today we learned...

- 🚂 What Ray Serve is and how it works
- 🔬 How to use Serve for production services
- 🔍 Why to choose Serve for AI-based apps

More Resources

For further exploration with Ray, Anyscale, and LLMs.

Reading list.

Ray Education GitHub

Access bonus notebooks and scripts about Ray.

Ray documentation

API references and user guides.

Anyscale Blogs

Real world use cases and announcements.

YouTube Tutorials

Video walkthroughs about learning LLMs with Ray.

Connect with the community.

👋 Join the community

Attend events, *subscribe to newsletter*, *follow on Twitter*.

🙋 Get support

Join Ray Slack, *ask questions on forum*, *open an issue*.

🔭 Contribute to Ray

Read contributor guide, *create a pull request*.

Thank you!

We hope to meet again.